


STAT6181

Computational Statistics I

Lecture 2



Optimization



Overview

- R review
- Finding the roots of an equation
- Optimizing functions



Introduction to R

R free software:

- We can run simulations
- We code 'new' statistical methods
- R is not particularly good at heavy computational problems
- R's memory management is not the best in the world.



R – the basics (it's a calculator)

```
> 1+1
```

```
[1] 2
```

```
>
```

```
> 3*3
```

```
[1] 9
```

```
>
```

```
> 4^3
```

```
[1] 64
```

```
>
```

```
> 100/4
```

```
[1] 25
```

```
>
```

```
>
```

```
> 2*sin(60)           # R uses the default Radians!
```

```
[1] -0.6096212
```

```
>
```

```
> log(100)           #To the base e
```

```
[1] 4.60517
```

```
>
```

```
> log10(100)         #To the base 10
```

```
[1] 2
```

```
>
```

R- creation of columns



```
>x = c(1,2,3, 4, 5, 6)
```

```
>y = c(7,8,9,10,11,12)
```

```
>z = c(1,2,3)
```

```
> x + y
```

```
[1] 8 10 12 14 16 18
```

```
> x+z
```

```
[1] 2 4 6 5 7 9
```

```
> x*y
```

```
[1] 7 16 27 40 55 72
```

```
> sum(x*y)
```

```
[1] 217
```

```
> x*z
```

```
[1] 1 4 9 4 10 18
```

```
> y^x
```

```
[1] 7 64 729 10000 161051 2985984
```

```
> y-z
```

```
[1] 6 6 6 9 9 9
```



R- Creation of Matrices and columns

```
> a = matrix ( c( 1,2,3,4), 2,2)
```

```
> a
```

```
      [,1] [,2]
```

```
[1,]  1  3
```

```
[2,]  2  4
```

```
> b = matrix ( c(1,2,3,4,5,6), 2,3)
```

```
> b
```

```
      [,1] [,2] [,3]
```

```
[1,]  1  3  5
```

```
[2,]  2  4  6
```

```
> c = c( 1,2)
```

```
> c
```

```
[1] 1 2
```

```
> d = c(1,2,3)
```

```
> d
```

```
[1] 1 2 3
```

```
>
```



R- Matrix mathematical manipulations

```
> a*a
```

```
      [,1] [,2]
```

```
[1,]    1    9
```

```
[2,]    4   16
```

```
> a*b
```

Error in a * b : non-conformable arrays

```
> a%%a
```

```
      [,1] [,2]
```

```
[1,]    7   15
```

```
[2,]   10   22
```

```
> a%%b
```

```
      [,1] [,2] [,3]
```

```
[1,]    7   15   23
```

```
[2,]   10   22   34
```

```
> a%%t(b)
```

Error in a %% t(b) : non-conformable arguments

```
> a%%c
```

```
      [,1]
```

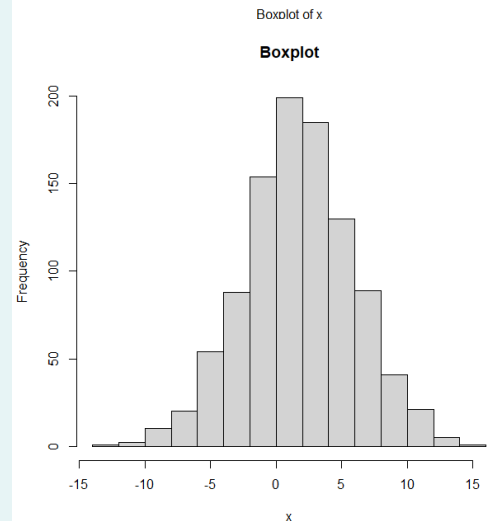

```
[1,]    7
```

```
[2,]   10
```


Plotting data – univariate data

- There are many packages that exist that can assist with graphical plots (eg ggplot etc.)
- I will go through the basics:
Univariate continuous data:

```
> x = rnorm(1000, 2,4)
> boxplot(x, col='red', xlab = 'Boxplot of x',ylab = 'x')
> hist(x, main = "Boxplot", xlab = 'x', ylab = 'Frequency')
```



Lets just look at the distributions a bit and some R functions



These are some of the common function associated with distributions found in R.

`rnorm`

`pnorm`

`dnorm`

`qnorm`

These functions enable you to compute probabilities find generate random data from almost **any** of our 'common' distributions.



Some other important R functions

```
> # This is the same as computing the pdf of the normal with x = 0, mu = 0 and  
> # sigma = 0. The dnorm function takes three main arguments. In this example, we are  
> #calculating the “density” at x=0 from a Normal( 2, 16) distribution.
```

```
>
```

```
> dnorm(0, mean = 2, sd = 4)
```

```
[1] 0.08801633
```

```
> #The function pnorm returns the integral from  $-\infty$  to q of the pdf of the normal  
> #distribution where q is a Z-score. Try to guess the value of pnorm(0,0,1).  
> #pnorm has the same default mean and sd arguments as dnorm).
```

```
> # pnorm gives  $P(X \leq x)$ 
```

```
>
```

```
> pnorm(2, mean = 2, sd = 4)
```

```
[1] 0.5
```

```
>
```

```
> pnorm(3, mean = 2, sd = 4)
```

```
[1] 0.5987063
```

```
>
```



Some more functions - qnorm

> #The **qnorm** function is simply the inverse of the cdf. So it like the inverse of pnorm.

> #You can use qnorm to determine the answer to the question: What is the

> #qnorm answers the question - what is the Z-score of the pth quantile of the normal distribution?

>

> # What is the Z-score of the 50th quantile of the normal distribution?

> **qnorm(.5)**

[1] 0

>

> # What is the Z-score of the 96th quantile of the normal distribution?

> qnorm(.96)

[1] 1.750686

>

> # What is the Z-score of the 46th quantile of the normal distribution with mean 2 and variance 16?

> qnorm(.46, 2, 4)

[1] 1.598265

>



rnorm revisited

```
> #rnorm revisited
```

```
> #If you want to generate a vector of normally distributed random numbers, > #rnorm is the function you should use. The first argument n is the number of numbers you want to generate, followed by the mean and sd arguments.
```

```
> rnorm(5, 2, 4)
```

```
[1] 5.114324 6.290950 3.102321 2.846167 4.912762
```

```
>
```

```
> rnorm(5, 2, 4)
```

```
[1] -2.3878274 -1.3153996 2.4711667 1.5269966 -0.4817649
```

```
>
```

```
> #So notice that I ran these twice and I got 5 different numbers. I need set.seed if I want the same random numbers.
```

```
> # For example,
```

```
>
```

```
> set.seed(10)
```

```
> rnorm(5, 3, 4)
```

```
[1] 3.0749847 2.2629898 -2.4853222 0.6033291 4.1781805
```

```
>
```

```
> set.seed(15)
```

```
> rnorm(5, 3, 4)
```

```
[1] 4.035291 10.324483 1.641526 6.588793 4.952065
```

```
>
```

```
> #Try set.seed = 10 now and see what happens
```

```
> set.seed(10)
```

```
> rnorm(5, 3, 4)
```

```
[1] 3.0749847 2.2629898 -2.4853222 0.6033291 4.1781805
```



Functions in R

- A function is a “set” of code that takes in objects/values/etc and returns a ‘set’ of outputs.
- The inputs can be objects in R.
- The outputs can be objects in R.




A simple R function

```
sumtwocolumns <- function( x, y)  
{  
  z = x+y  
  return(z)  
}
```

Lets go through line by line!!!

Now, how do you run the function



- **Step 1:** You need to highlight the code for the function and run it
- **Step 2:** You have the objects/inputs you need
- **Step 3:** Create a single line to run the function

#You need two input columns x and y, so I will create two columns x and y

```
x = c(1,2,3,4)
```

```
y = c(5,6,7,8)
```



Then, I run the function

```
>sumtwocolumns(x,y)
```


```
[1] 6 8 10 12
```

```
>k = sumtwocolumns(x,y)
```

```
>k
```

```
[1] 6 8 10 12
```


Can you see what is happening here?



```
sum.of.squares <- function(x,y)
{
x^2 + y^2
}
```

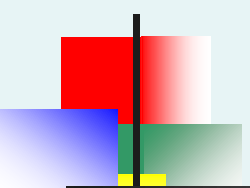
```
> d = sum.of.squares(x,y)
```

```
> d
```

```
[1] 26  40  58  80
```

```
>
```

Functions can return more than one item...an example



```
two.returns <- function(x,y)
{
  z=x^2 + y^2
  return(list(z,min(x), max(y)))
}
```

This function returns 3 items


```
> two.returns(x,y)
[[1]]
[1] 26 40 58 80
```

Now lets call the function

```
[[2]]
[1] 1
```

```
[[3]]
[1] 8
```

Is there an easier way to call the function?



```
> m = two.returns(x,y)
```

```
> m[[1]]
```

```
[1] 26 40 58 80
```

```
> m[[2]]
```

```
[1] 1
```

```
> m[[3]]
```

```
[1] 8
```

```
> m[[4]]
```

```
Error in m[[4]] : subscript out of bounds
```



Can R differentiate?

YES!! We can use the D() function

Example:

```
> f = expression(-6*x^3 - sqrt(x))
```

Define expression

```
> D(f,'x')
```

Find first derivative

```
-(6 * (3 * x^2) + 0.5 * x^-0.5)
```

```
> D(D(f,'x'), 'x')
```

Find second derivative

```
-(6 * (3 * (2 * x)) + 0.5 * (-0.5 * x^-1.5))
```



Can R integrate?


```
> integrand <- function(x) {-6*x^3 - sqrt(x)}  
>  
> integrate(integrand, lower = .1, upper = .25)  
-0.06796086 with absolute error < 7.5e-16
```

Roots of an equation – Plotting the function...

It is always good if we can get a graph of the curve In the code that follows, we try to get a graph of the function $f(x) = 1/2 - \exp\{-(x^2)\}$

```
f <- function(x) .5-exp(-(x^2))
v <- seq(0,2,length=1000)
plt <- function(f,a0,b0)
{
  plot(v,f(v),type="l", ylim=c(-.5,.5), xlim=c(0,2),
    main=sprintf("Interval length = %f", (b0-a0)))
  abline(h=0)
  segments(a0,.1,a0,-.1,col=2,lty=1)
  points(a0,.1,col=2)
  points(a0,-.1,col=2)
  segments(b0,.1,b0,-.1,col=2,lty=1)
  points(b0,.1,col=2)
  points(b0,-.1,col=2)
  segments(mean(c(a0,b0)), .1, mean(c(a0,b0)), -.1, col=4, lty=1)
  points(mean(c(a0,b0)),.1,col=4)
  points(mean(c(a0,b0)),-.1,col=4)
}
iter1 <- function(f,l)
{
  m = mean(l)
  if( f(l[1])*f(m) < 0 ) return( c(l[1],m) ) else return( c(m, l[2]) )
}
l <- c(0,2)
f <- function(x) .5-exp(-(x^2))
a0 <- l[1]; b0 <- l[2];
plt(f, a0, b0)
l = iter1(f,l)
```

...The graph that is produced...





The Bisection Method

This is the bisection function. `bisect` is a function that takes a function `f` and an interval (a,b) that contains the root followed by a tolerance level.



The Bisection Method

The bisection method requires you to provide a bracket, $c(a_0, b_0)$ as a starting point. The method proceeds by splitting the interval in half, and throws out the interval where the sign did not change, and repeats. Once the interval shrinks below a certain length, δ , it is considered to have converged. Put formally, at step k ,

- Calculate $m = (a_{k-1} + b_{k-1})/2$, the midpoint of the current bracket
- if $f(a_{k-1}) \cdot f(m) < 0$ (i.e. the sign does changes over the interval (a_{k-1}, m)), then $a_k = a_{k-1}, b_k = m$.
- otherwise $a_k = m, b_k = b_{k-1}$
- Repeat until $b_k - a_k < \delta$ for some small δ chosen beforehand



The Bisection Method

This is the bisection function. `bisect` is a function that takes a function `f` and an interval (a,b) that contains the root followed by a tolerance level.



The Bisection Function

```
bisect <- function(f, a, b, tol)
{
  l <- c(a,b)
  L <- l[2]-l[1]
  while( L > tol )
  {
    m <- mean(l)
    if( f(m)*f(l[1]) < 0 ) l = c(l[1],m) else l = c(m,l[2])
    L <- l[2]-l[1]
  }
  return(mean(l))
}
```



Running the bisection function

Define the function

```
fn <- function(x) 0.5-exp(-(x^2))
```

Use the bisection rule to get an estimate of the root. We need to supply

- the function,
- the interval that contains the root and a guess (starting value).
- The tolerance level.
- The bisection function returns the xvalue.

```
> bisect(fn, 0, 2, 1e-6)
```

```
[1] 0.8325543
```

We can get the y value by substituting this xvalue in the function

```
> fn(.8325543)
```

```
[1] -2.590558e-07
```



Example 2

In this example, we first generate a random sample of data from a normal distribution with mean 0 and variance 1. We will assume the variance is known and the mean is unknown. So there is just one parameter to estimate. This is how we generate the data:

```
X <- rnorm(100)
```

```
## We need the function. Can you see how I got the function below?
```

```
fn <- function(t) sum(X-t)
```

```
## Find an interval containing the root. We need this interval for the bisection function.
```

```
> c( fn(-1), fn(1) )
```

```
[1] 95.17979 -104.82021
```

```
## Apply the bisection function with tolerance level.
```

```
bisect(fn, -1, 1, 1e-7)
```

```
[1] -0.04820213
```

```
### Recall we can check our answer with the actual MLE which is the sample mean.
```

```
> mean(X)
```

```
[1] -0.0482021
```

So how close are we to what we had expected?



Newton Raphson Method

An extremely fast root-finding approach is Newton's method. This approach is also referred to as *Newton–Raphson iteration*, especially in univariate applications. Suppose that g' is continuously differentiable and that $g''(x^*) \neq 0$. At iteration t , the approach approximates $g'(x^*)$ by the linear Taylor series expansion:

$$0 = g'(x^*) \approx g'(x^{(t)}) + (x^* - x^{(t)})g''(x^{(t)}). \quad (2.8)$$

Since g' is approximated by its tangent line at $x^{(t)}$, it seems sensible to approximate the root of g' by the root of the tangent line. Thus, solving for x^* above, we obtain

$$x^* = x^{(t)} - \frac{g'(x^{(t)})}{g''(x^{(t)})} = x^{(t)} + h^{(t)}. \quad (2.9)$$

This equation describes an approximation to x^* that depends on the current guess $x^{(t)}$ and a refinement $h^{(t)}$. Iterating this strategy yields the updating equation for Newton's method:

$$x^{(t+1)} = x^{(t)} + h^{(t)}, \quad (2.10)$$



Newton Raphson continued

where $h^{(t)} = -g'(x^{(t)}) / g''(x^{(t)})$. The same update can be motivated by analytically solving for the maximum of the quadratic Taylor series approximation to $g(x^*)$, namely $g(x^{(t)}) + (x^* - x^{(t)})g'(x^{(t)}) + (x^* - x^{(t)})^2 g''(x^{(t)})/2$. When the optimization of g corresponds to an MLE problem where $\hat{\theta}$ is a solution to $l'(\theta) = 0$, the updating equation for Newton's method is

$$\theta^{(t+1)} = \theta^{(t)} - \frac{l'(\theta^{(t)})}{l''(\theta^{(t)})}. \quad (2.11)$$

Newton Raphson Method


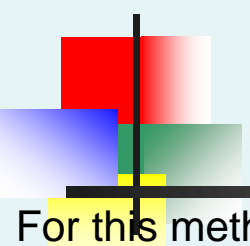


FIGURE 2.3 Illustration of Newton's method applied in Example 2.2. At the first step, Newton's method approximates g' by its tangent line at $x^{(0)}$, whose root $x^{(1)}$ serves as the next approximation of the true root x^* . The next step similarly yields $x^{(2)}$, which is already quite close to x^* .

A simple Newton-Raphson R function Method



For this method, we need the function fast well as the first derivative df.

#This is the function

```
f <- function(x) .5-exp(-(x^2))
```

#This is the derivative of the function

```
df <- function(x) 2*x*exp(-(x^2))
```

Lets plot the graph first

```
v <- seq(0,2,length=1000)
```

```
plot(v,f(v),type="l",col=8)
```

```
abline(h=0)
```

```
x0 <- 1.5
```

```
segments(x0,0,x0,f(x0),col=2,lty=3)
```


```
slope <- df(x0)
```

```
g <- function(x) f(x0) + slope*(x - x0)
```

```
v = seq(x0 - f(x0)/df(x0),x0,length=1000)
```

```
lines(v,g(v),lty=3,col=4)
```

The plot of the graph



Lets do one iteration of the Newton Raphson before proceeding...



##Starting point

$x_0 = 1.5$

This is our NR update

$x_0 \leftarrow x_0 - f(x_0)/df(x_0)$

```
> x0 <- x0 - f(x0)/df(x0)
```

```
>
```

```
> x0
```


```
[1] 0.252044
```

Now lets use an NR function

```
newton <- function(x0, f, df, d2f, tol=1e-4, pr=FALSE)
{
  k <- 1
  fval <- f(x0)
  grad <- df(x0)
  hess <- d2f(x0)
  xk_1 <- x0
  cond1 <- sqrt( sum(grad^2) )
  cond2 <- Inf
  if( (cond1 < tol) ) return(x0)
  while( (cond1 > tol) & (cond2 > tol) )
  {
    L <- 1
    bool <- TRUE
    while(bool == TRUE)
    {
      xk <- xk_1 - L * solve(hess) %*% grad
      if( f(xk) > fval )
      {
        bool = FALSE
        grad <- df(xk)
        fval <- f(xk)
        hess <- d2f(xk)
      } else
      {
        L = L/2
        if( abs(L) < 1e-20 ) return("Failed to find uphill step - try new start values")
      }
    }
    cond1 <- sqrt( sum(grad^2) )
    cond2 <- sqrt( sum( (xk-xk_1)^2 ) )/(tol + sqrt(sum(xk^2)))
    k <- k + 1
    xk_1 <- xk
  }

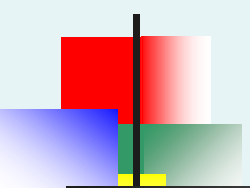
  if(pr == TRUE) print( sprintf("Took %i iterations", k) )
  return(xk)
}
```

Illustrate the function using a simple function



```
> f <- function(x) exp(-(x^2))  
> df <- function(x) -2*x*f(x)  
> d2f <- function(x) -2*f(x)-2*x*df(x)  
> newton(2/3, f, df, d2f, 1e-7)  
      [,1]  
[1,] -3.233794e-09
```

Now illustrate the NR method with two parameters




This is the approach, we first simulate 100 values from a known distribution with the parameters KNOWN (off course).

We will generate $n=100$ values from a $\text{Normal}(2,4)$. Then we will try to use the NR method to estimate those parameters.

```
X <- rnorm(100, mean=2, sd=sqrt(4))  
n <- 100
```

- We will need define the likelihood function
- Then we take the logs of that function
- We will see why it is the sum of logs
- It requires a vector $t = c(\mu, \sigma)$

Define the sum of the likelihood function




```
f <- function(t)
{
  if( t[2] > 0)
  {
    return( sum( dnorm(X, mean=t[1], sd=sqrt(t[2]), log=TRUE) ) )
  } else
  {
    return(-Inf)
  }
}
```

The vector of score functions – first differentials



```
df <- function(t)
{
mu <- t[1]
sig <- t[2]
g <- rep(0,2)
g[1] <- (1/sig) * sum(X - mu)
g[2] <- (-n/(2*sig)) + sum( (X-mu)^2 )/(2*sig^2)
return(g)
}
```


The Hessian matrix of second differentials



```
d2f <- function(t)
{
  mu <- t[1]; sig <- t[2];
  h <- matrix(0,2,2)
  h[1,1] <- -n/sig
  h[2,2] <- (n/(2*sig^2)) - sum( (X-mu)^2 )/(sig^3)
  h[1,2] <- -sum( (X-mu) )/(sig^2)
  h[2,1] <- h[1,2]
  return(h)
}
```



Now apply the NR

We simply apply newton's method now.

```
> newton( c(0,1), f, df, d2f)
```

```
      [,1]
```

```
[1,]  2.050354
```

```
[2,]  3.151090
```

Lets compare the results with actual MLEs

```
> c( mean(X), (n-1)*var(X)/n)
```

```
[1] 2.050354    3.151090
```